

```
In [ ]: import numpy as np
        from sklearn.datasets import fetch_openml
        from sklearn.decomposition import PCA

        # Load the MNIST dataset
        mnist = fetch_openml('mnist_784', version=1)
        X, y = mnist['data'], mnist['target']

        # Split the dataset into training and test sets
        X_train, X_test = X[60000:], X[60000:]
        y_train, y_test = y[60000:], y[60000:]

        # Apply PCA with 20 principal components
        pca = PCA(n_components=20)
        X_train_pca = pca.fit_transform(X_train)

        # Print the top 20 eigenvalues
        print("Top 20 Eigenvalues:")
        print(pca.explained_variance_)

        # Print the explained variance ratios of the principal components
        print("\nExplained_Variance_Ratio:")
        print(pca.explained_variance_ratio_)

        #user/local/lib/python3.8/dist-packages/sklearn/datasets/_openml.py:968: FutureWarning: The default value of 'parser'
        #will change from 'auto' to 'auto-raw' in 1.4. You can set the parser to 'auto' to silence this warning. Therefore, an
        #error/factor will be raised from 3.4 if the dataset is dense and pandas is not installed. Note that the pandas parser
        #may return different data types. See the Notes Section in fetch_openml's API doc for details.
        warn(

        Top 20 Eigenvalues:
        [332724.68744657 243283.9390765 211567.36705827 184776.38586212
        166926.8313063 147844.8617525 121278.20267351 98074.42953629
        94096.24875017 80809.63240654 72133.61931102 60353.29678698
        55826.7202427 58913.6749777 51243.34856213 50841.7032562
        45405.20654754 43777.97588424 40761.02970332 39518.12208522]

        Explained_Variance_Ratio:
        [0.89704684 0.87095924 0.86160689 0.85380419 0.84687897 0.84122121
        0.8272193 0.82683895 0.82762029 0.82357901 0.8210919 0.82022991
        0.8175814 0.816921 0.81578629 0.81482913 0.81234345 0.81276883
        0.81187137 0.81152638]
```

```

from sklearn.datasets import fetch_opheml
from sklearn.model_selection import train_test_split
from sklearn.decomposition import PCA
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Load the MNIST dataset
mnist = fetch_opheml('mnist_784', version=1)
X, y = mnist['data'], mnist['target']

# Split the dataset into training, validation, and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

X_train_val, y_train_val, X_test_val, y_test_val = train_test_split(X_train_val, y_train_val, test_size=0.25, random_state=42) # 0.25
# 0.8 = 0.2

# Apply PCA with 28 principal components
pca = PCA(n_components=28)
X_train_pca = pca.fit_transform(X_train)
X_val_pca = pca.transform(X_val)
X_test_pca = pca.transform(X_test)

# Train logistic regression model with increased max_iter
log_reg = LogisticRegression(max_iter=10000) # Increase max_iter value
log_reg.fit(X_train_pca, y_train)

# Calculate training, validation, and test accuracy
train_accuracy = accuracy_score(y_train, log_reg.predict(X_train_pca))
val_accuracy = accuracy_score(y_val, log_reg.predict(X_val_pca))
test_accuracy = accuracy_score(y_test, log_reg.predict(X_test_pca))

print("Logistic Regression Model's Accuracy:")
print("Training Accuracy:", train_accuracy)
print("Validation Accuracy:", val_accuracy)
print("Test Accuracy:", test_accuracy)

# User/library/python path/directories/sklearn/datasets/opheml.py:968: FutureWarning: The default value of 'parser'
# will change from "liac-arff" to "auto" in 1.4. You can set 'parser='auto"' to silence this warning. Therefore, an
# "ImportError" will be raised from 3.4 if the dataset is dense and pandas is not installed. Note that the pandas parser
# may return different data types. See the Notes Section in fense and pandas is not installed. Note that the pandas parser
# warn()

Logistic Regression Model's Accuracy:
Training Accuracy: 0.878952809953899
Validation Accuracy: 0.8728
Test Accuracy: 0.877874248574428

In [ ]:
import numpy as np
from sklearn.datasets import fetch_opheml
from sklearn.model_selection import train_test_split
from sklearn.decomposition import PCA
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
from sklearn.decomposition import TruncatedSVD
from sklearn.linear_model import SGDClassifier
from sklearn.preprocessing import StandardScaler

# Load the MNIST dataset
mnist = fetch_opheml('mnist_784', version=1)
X, y = mnist['data'], mnist['target']

# Split the dataset into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Define pipeline with PCA and logistic regression
pipe = Pipeline([
    ('reduce_dim', PCA(n_components=10)), # Use PCA for dimensionality reduction
    ('clf', LogisticRegression()) # Use LogisticRegression for Logistic regression
])

# Define parameter grid for grid search
param_grid = {
    'reduce_dim__n_components': [1, 5, 10, 15, 20], # Use a smaller range for the number of components
    'clf__max_iter': [1000, 2000], # Increase max_iter if necessary
}

# Perform grid search with cross-validation
grid_search = GridSearchCV(pipe, param_grid, cv=3, n_jobs=-1) # Use parallel processing
grid_search.fit(X_train, y_train)

```

```

# Report the number of principal components that achieve the highest validation accuracy
best_num_components = grid_search.best_params_['reduce_dim_n_components']
print("Number of principal components with highest Validation Accuracy: ", best_num_components)

# Calculate training and test accuracy using the selected number of principal components
train_accuracy = accuracy_score(y_train, grid_search.predict(X_train))
test_accuracy = accuracy_score(y_test, grid_search.predict(X_test))

# Report training and test accuracy
print("Training Accuracy with Selected Number of Principal Components: ", train_accuracy)
print("Test Accuracy with Selected Number of Principal Components: ", test_accuracy)

# /usr/local/lib/python3.8/dist-packages/sklearn/datasets/openml.py:968: FutureWarning: The default value of 'parser'
# will change from 'lisc-ascii' to 'auto' in 1.4. You can set 'parser=auto': to silence this warning. Therefore, an
# 'error' will be raised from 1.4 if the dataset is dense and pandas is not installed. Note that the pandas para
# r may return different data types. See the Notes section in fetch_openml's API doc for details.
# warn()

Number of Principal Components with Highest Validation Accuracy: 28
Training Accuracy with Selected Number of Principal Components: 0.874785142857143
Test Accuracy with Selected Number of Principal Components: 0.8775428571428571

In [ ]:
import pandas as pd
from sklearn.model_selection import train_test_split

# Load the College dataset
college_data = pd.read_csv("college_data.csv") # Assuming you have the dataset saved as college_data.csv

# Split the data into features (X) and target variable (y)
X = college_data.drop("AcceptanceRate", axis=1)
y = college_data["AcceptanceRate"]

# Split the data into training and test sets with 80% training and 20% test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Print the shape of the training and test sets
print("Training set shape: ", X_train.shape, y_train.shape)
print("Test set shape: ", X_test.shape, y_test.shape)

```

```
[ 1] import pandas as pd
[ 2] from sklearn.cross_decomposition import PLSRegression
[ 3] from sklearn.model_selection import train_test_split, GridSearchCV
[ 4] from sklearn.pipeline import Pipeline
[ 5] from sklearn.preprocessing import StandardScaler
[ 6] from sklearn.metrics import mean_squared_error
[ 7]
[ 8] # Load the College dataset
[ 9] college_data = pd.read_csv("college_data.csv") # Assuming you have the dataset saved as college_data.csv
[10]
[11] # Split the data into features (X) and target variable (y)
[12] X = college_data.drop(columns=['Accept', 'Apps', 'AcceptanceRate']) # Exclude 'Accept' and 'Apps' as requested
[13] y = college_data['AcceptanceRate']
[14]
[15] # Convert the target variable to numeric
[16] y = pd.to_numeric(y, errors='coerce') # Convert any non-numeric values to NaN
[17]
[18] # Drop rows with NaN values in the target variable
[19] X = X.dropna(subset=['AcceptanceRate'])
[20] y = y.dropna()
[21]
[22] # Split the data into training and test sets with 80% training and 20% test
[23] X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Define the PLS regression pipeline
pipeline = Pipeline([
    ('scaler', StandardScaler()), # Scale the features
    ('pls', PLSRegression()), # PLS regression model
])

# Define the parameter grid for grid search
param_grid = {
    'pls_n_components': range(1, min(X_train.shape[1], 10)), # M: Number of components for PLS
}

# Perform grid search with cross-validation
grid_search = GridSearchCV(pipeline, param_grid, cv=5, scoring='neg_mean_squared_error')
grid_search.fit(X_train, y_train)

# Get the best PLS model
best_pls_model = grid_search.best_estimator_

# Predict on the test set
y_pred = best_pls_model.predict(X_test)

# Calculate the test error (Mean Squared Error)
test_error = mean_squared_error(y_test, y_pred)
print("Test Error:", test_error)

# Report the value of M selected by cross-validation
best_M = grid_search.best_params_['pls_n_components']
print("Value of M selected by cross-validation:", best_M)

In [ ]: import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeRegressor, plot_tree
from sklearn.metrics import mean_squared_error

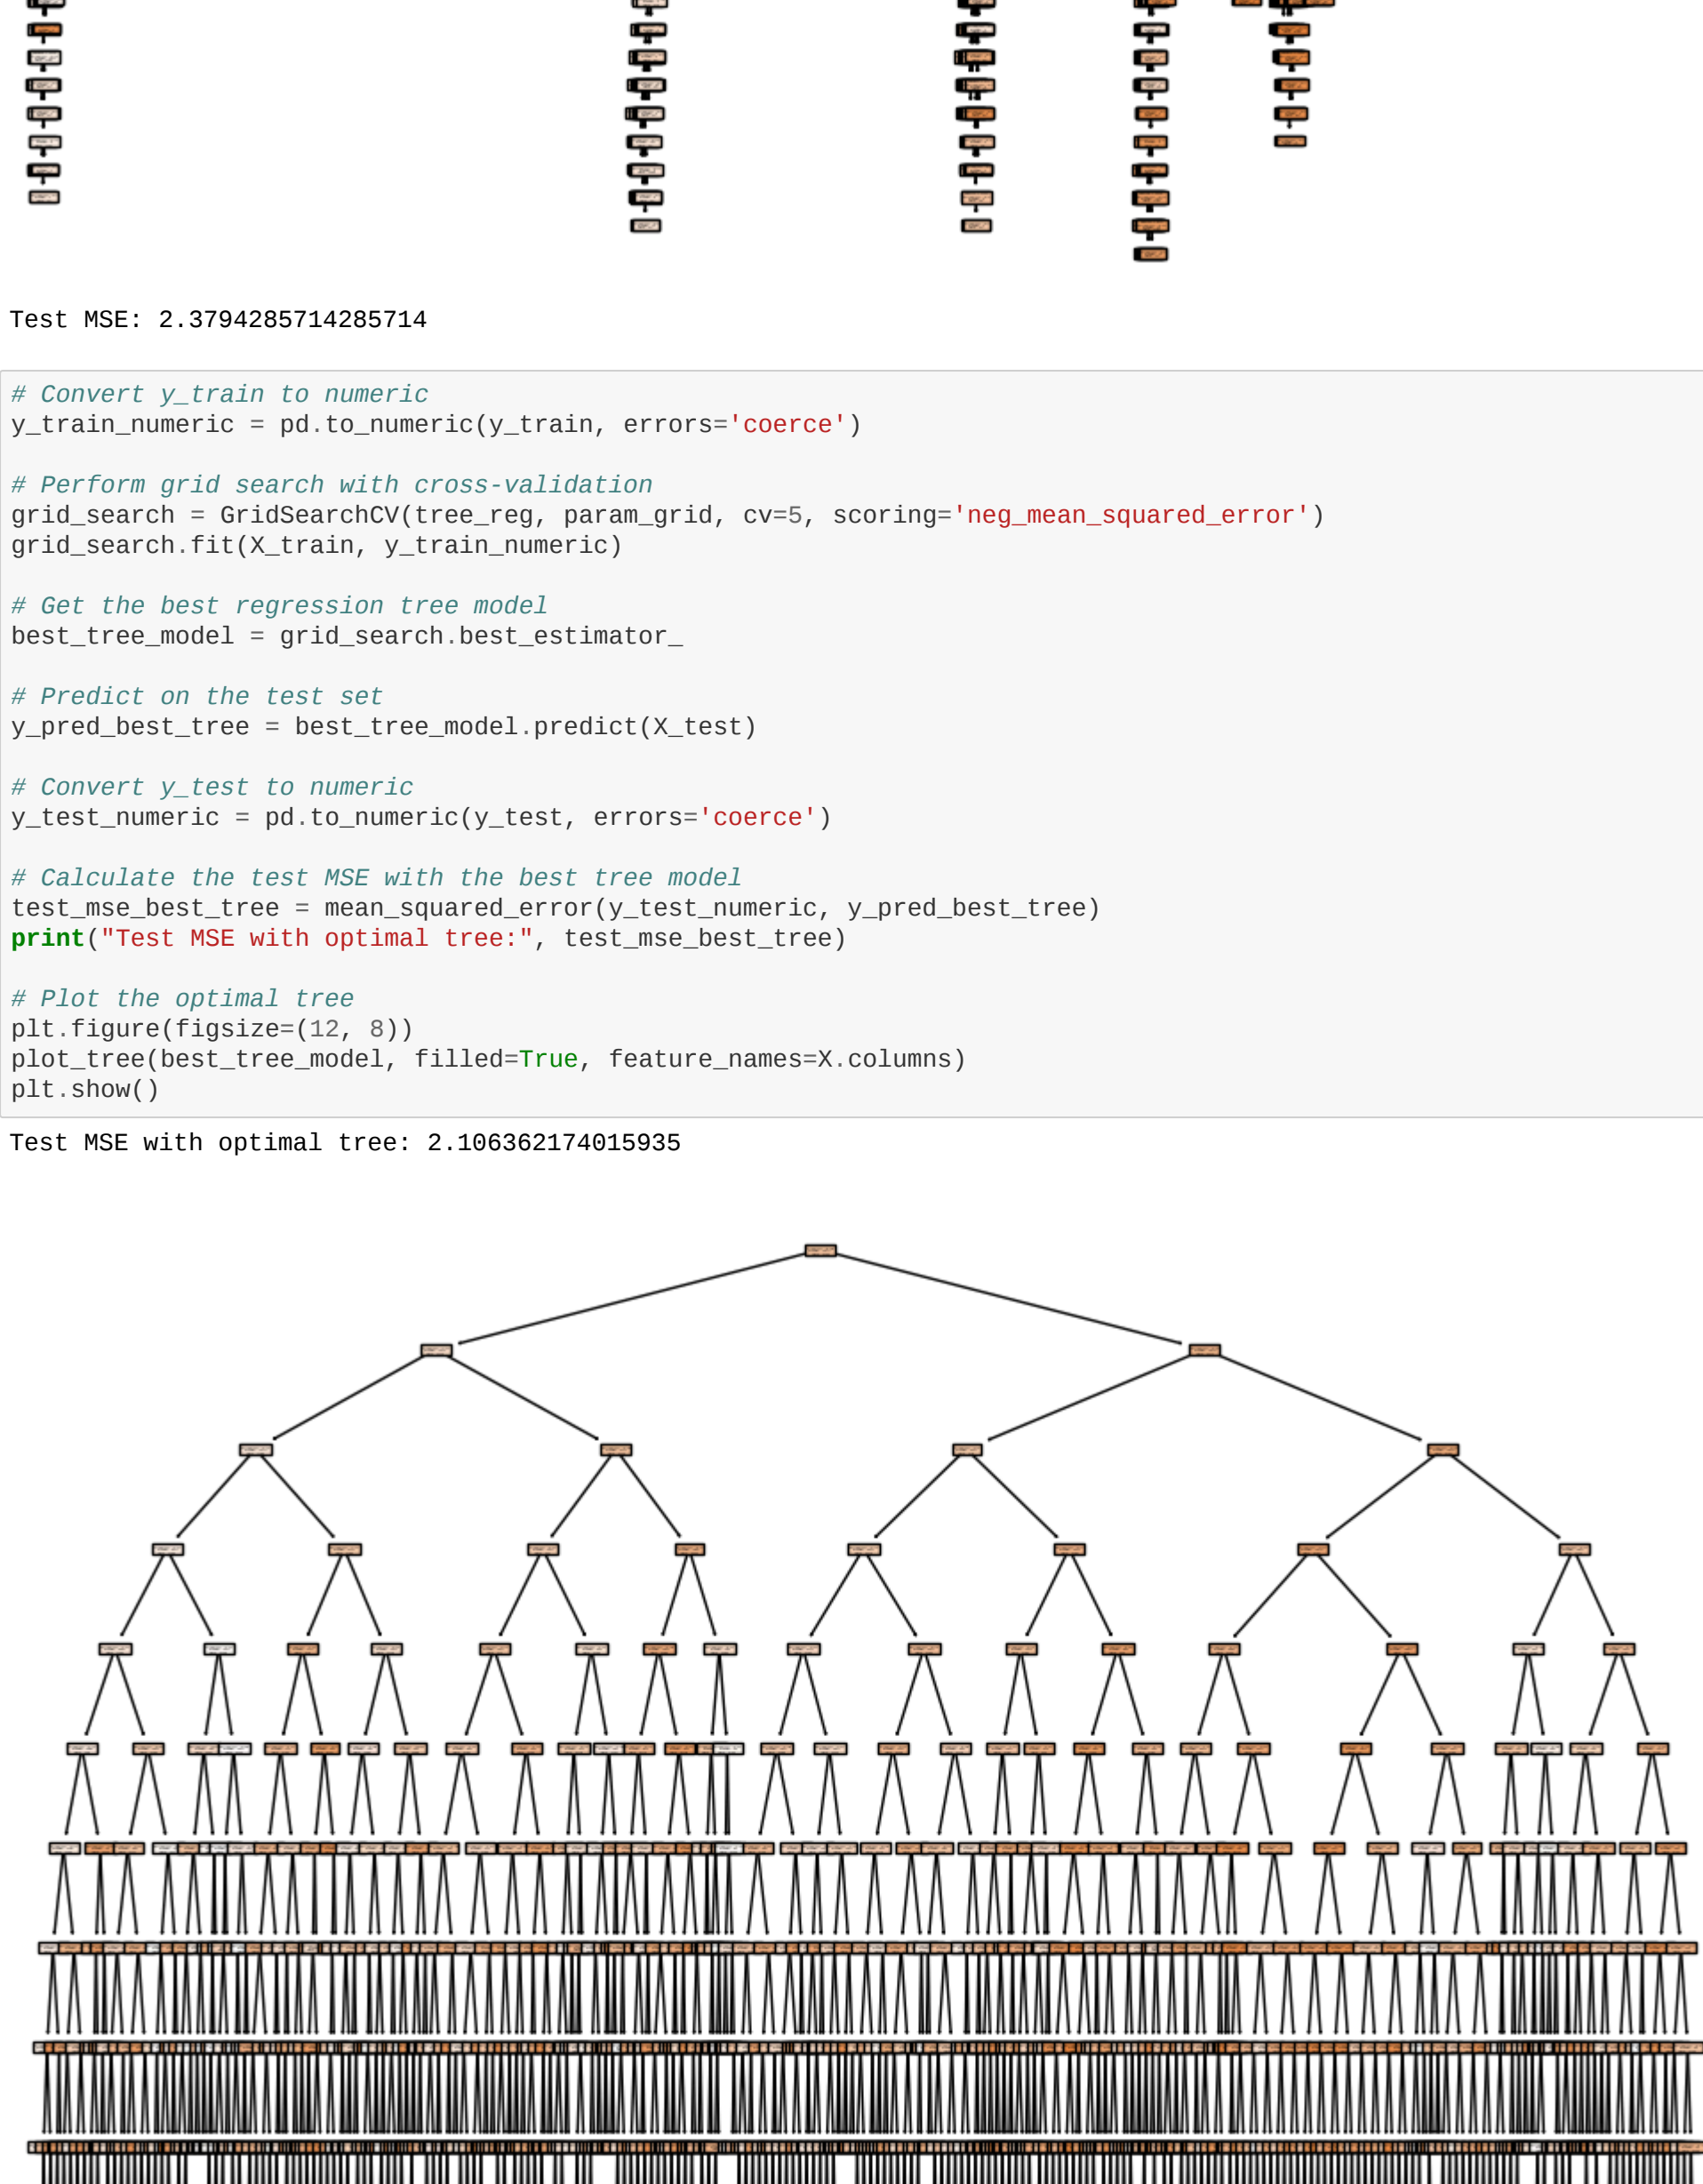
# Fit a regression tree to the training set
tree_reg = DecisionTreeRegressor(random_state=42)
tree_reg.fit(X_train, y_train)

# Plot the tree
plt.figure(figsize=(12, 8))
plot_tree(tree_reg, filled=True, feature_names=X.columns)
plt.show()

# Predict on the test set
y_pred_tree = tree_reg.predict(X_test)

# Convert y_test to numeric
y_test_numeric = pd.to_numeric(y_test, errors='coerce')

# Calculate the test MSE
test_mse_tree = mean_squared_error(y_test_numeric, y_pred_tree)
print("Test MSE:", test_mse_tree)
```



```
In [ ]: from sklearn.model_selection import train_test_split
        from sklearn.ensemble import BaggingRegressor
        from sklearn.tree import DecisionTreeRegressor
        from sklearn.metrics import mean_squared_error

        # Generate random feature matrix X and target variable y for testing
        np.random.seed(42)
        X = np.random.rand(100, 10) # Generating 100 samples with 10 features each
        y = np.random.rand(100) # Generating 100 random target values

        # Split the data into training and test sets
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

        # Create a Bagging Regressor with DecisionTreeRegressor as base estimator
        bagging_reg = BaggingRegressor(base_estimator=DecisionTreeRegressor(), n_estimators=100, random_state=42)

        # Fit the model on the training set
        bagging_reg.fit(X_train, y_train)
```

```
# Predict on the test set
y_pred_baggig = bagging_reg.predict(X_test)

# Calculate the test MSE
test_mse_baggig = mean_squared_error(y_test, y_pred_baggig)
print("Test MSE with Bagging Regressor:", test_mse_baggig)

# Get feature importances
feature_importances = np.mean([
    tree.feature_importances_ for tree in bagging_reg.estimators_
], axis=0)

# Print feature importances
print("Feature Importances:")
for i, importance in enumerate(feature_importances):
    print(f"Feature {i + 1}: {importance}")

#user/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:168: FutureWarning: 'base_estimator' was renamed to
'estimator' in version 1.2 and will be removed in 1.4.
warnings.warn(

Test MSE with Bagging Regressor: 0.18802761140790776
Feature Importances:
Feature 1: 0.6768873866652466
```

```

Feature 2: 0.153362414959867
Feature 3: 0.118412580743817
Feature 4: 0.67784038116289267
Feature 5: 0.07865517129463857
Feature 6: 0.6867755095189847
Feature 7: 0.1159942219607345
Feature 8: 0.07287481703840388
Feature 9: 0.1960412220414717
Feature 10: 0.11356981375682777

```

```

In [ ]: from sklearn.ensemble import RandomForestRegressor
        from sklearn.metrics import mean_squared_error

        # Generate random feature matrix X and target variable y for testing
        np.random.seed(42)
        X = np.random.random(100, 10) # Generating 100 samples with 10 features each
        y = np.random.random(100)     # Generating 100 random target values

        # Split the data into training and test sets
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

        # Define a range of values for the number of trees (n) and number of features considered at each split
        n_estimators_values = [10, 50, 100] # Number of trees
        max_features_values = ['sqrt', 'log2', 'log1p'] # Number of features considered at each split

```

```
# Iterate over different combinations of parameters and evaluate the performance
for estimators in n_estimators_values:
    for max_features in max_features_values:
        # Create a Random Forest regressor
        forest_reg = RandomForestRegressor(n_estimators=n_estimators, max_features=max_features, random_state=42)

        # Fit the model on the training set
        forest_reg.fit(X_train, y_train)

        # Calculate the test MSE
        y_pred_forest = forest_reg.predict(X_test)

        # Calculate the test MSE
        test_mse_forest = mean_squared_error(y_test, y_pred_forest)
        print(f'Test MSE with Random Forest (n_estimators={n_estimators}, max_features={max_features}): {test_mse_forest}')

# Import sklearn ensemble module
from sklearn.ensemble import RandomForestRegressor

# Import numpy module
import numpy as np

# Import pandas module
import pandas as pd
```

```

# Use this parameter as it is also the default value for RandomForestRegressor and ExtraTreeRegressor.
warn(
    "Warning: 'max_features' is deprecated in 1.1 and will be removed in 1.3. To keep the same behavior, explicitly set 'max_features='auto' or have 'warn' as this parameter as it is also the default value for RandomForestRegressor and ExtraTreeRegressor.",
    warn
)

Test MSE with Random Forest (n_estimators=10, max_features=auto): 0.09562456695990777
Test MSE with Random Forest (n_estimators=10, max_features=log2): 0.08422418025525121
Test MSE with Random Forest (n_estimators=10, max_features=sqrt): 0.08422418025525121
Test MSE with Random Forest (n_estimators=50, max_features=auto): 0.0839286268283444
Test MSE with Random Forest (n_estimators=50, max_features=log2): 0.0839286268283444
Test MSE with Random Forest (n_estimators=50, max_features=sqrt): 0.0839286268283444
Test MSE with Random Forest (n_estimators=100, max_features=auto): 0.10539911275895186

# /usr/local/lib/python3.6/dist-packages/sklearn/ensemble/_forest.py:413: FutureWarning: 'max_features='auto'' has been
# deprecated in 1.1 and will be removed in 1.3. To keep the same behavior, explicitly set 'max_features='auto' or have
# 'warn' as this parameter as it is also the default value for RandomForestRegressor and ExtraTreeRegressor.
warn(
    "Warning: 'max_features' is deprecated in 1.1 and will be removed in 1.3. To keep the same behavior, explicitly set 'max_features='auto' or have 'warn' as this parameter as it is also the default value for RandomForestRegressor and ExtraTreeRegressor.",
    warn
)

Test MSE with Random Forest (n_estimators=100, max_features=auto): 0.09461633000183335
Test MSE with Random Forest (n_estimators=100, max_features=log2): 0.09463339881833335

In [ ]: from sklearn.ensemble import GradientBoostingRegressor
        from sklearn.metrics import mean_squared_error

        # Generate random feature matrix x and target variable y for testing
        n, random_seed(42)

```

```
X = np.random.random(100, 10) # Generating 100 samples with 10 features each
y = np.random.random(100) # Generating 100 random target values

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create a Gradient Boosting regressor
gb_reg = GradientBoostingRegressor(random_state=42)

# Fit the model on the training set
gb_reg.fit(X_train, y_train)

# Predict on the test set
y_pred_gb = gb_reg.predict(X_test)

# Calculate the test MSE
test_mse_gb = mean_squared_error(y_test, y_pred_gb)
print("Test MSE with Gradient Boosting:", test_mse_gb)

# Plot the first tree in the ensemble
from sklearn.tree import plot_tree
import matplotlib.pyplot as plt

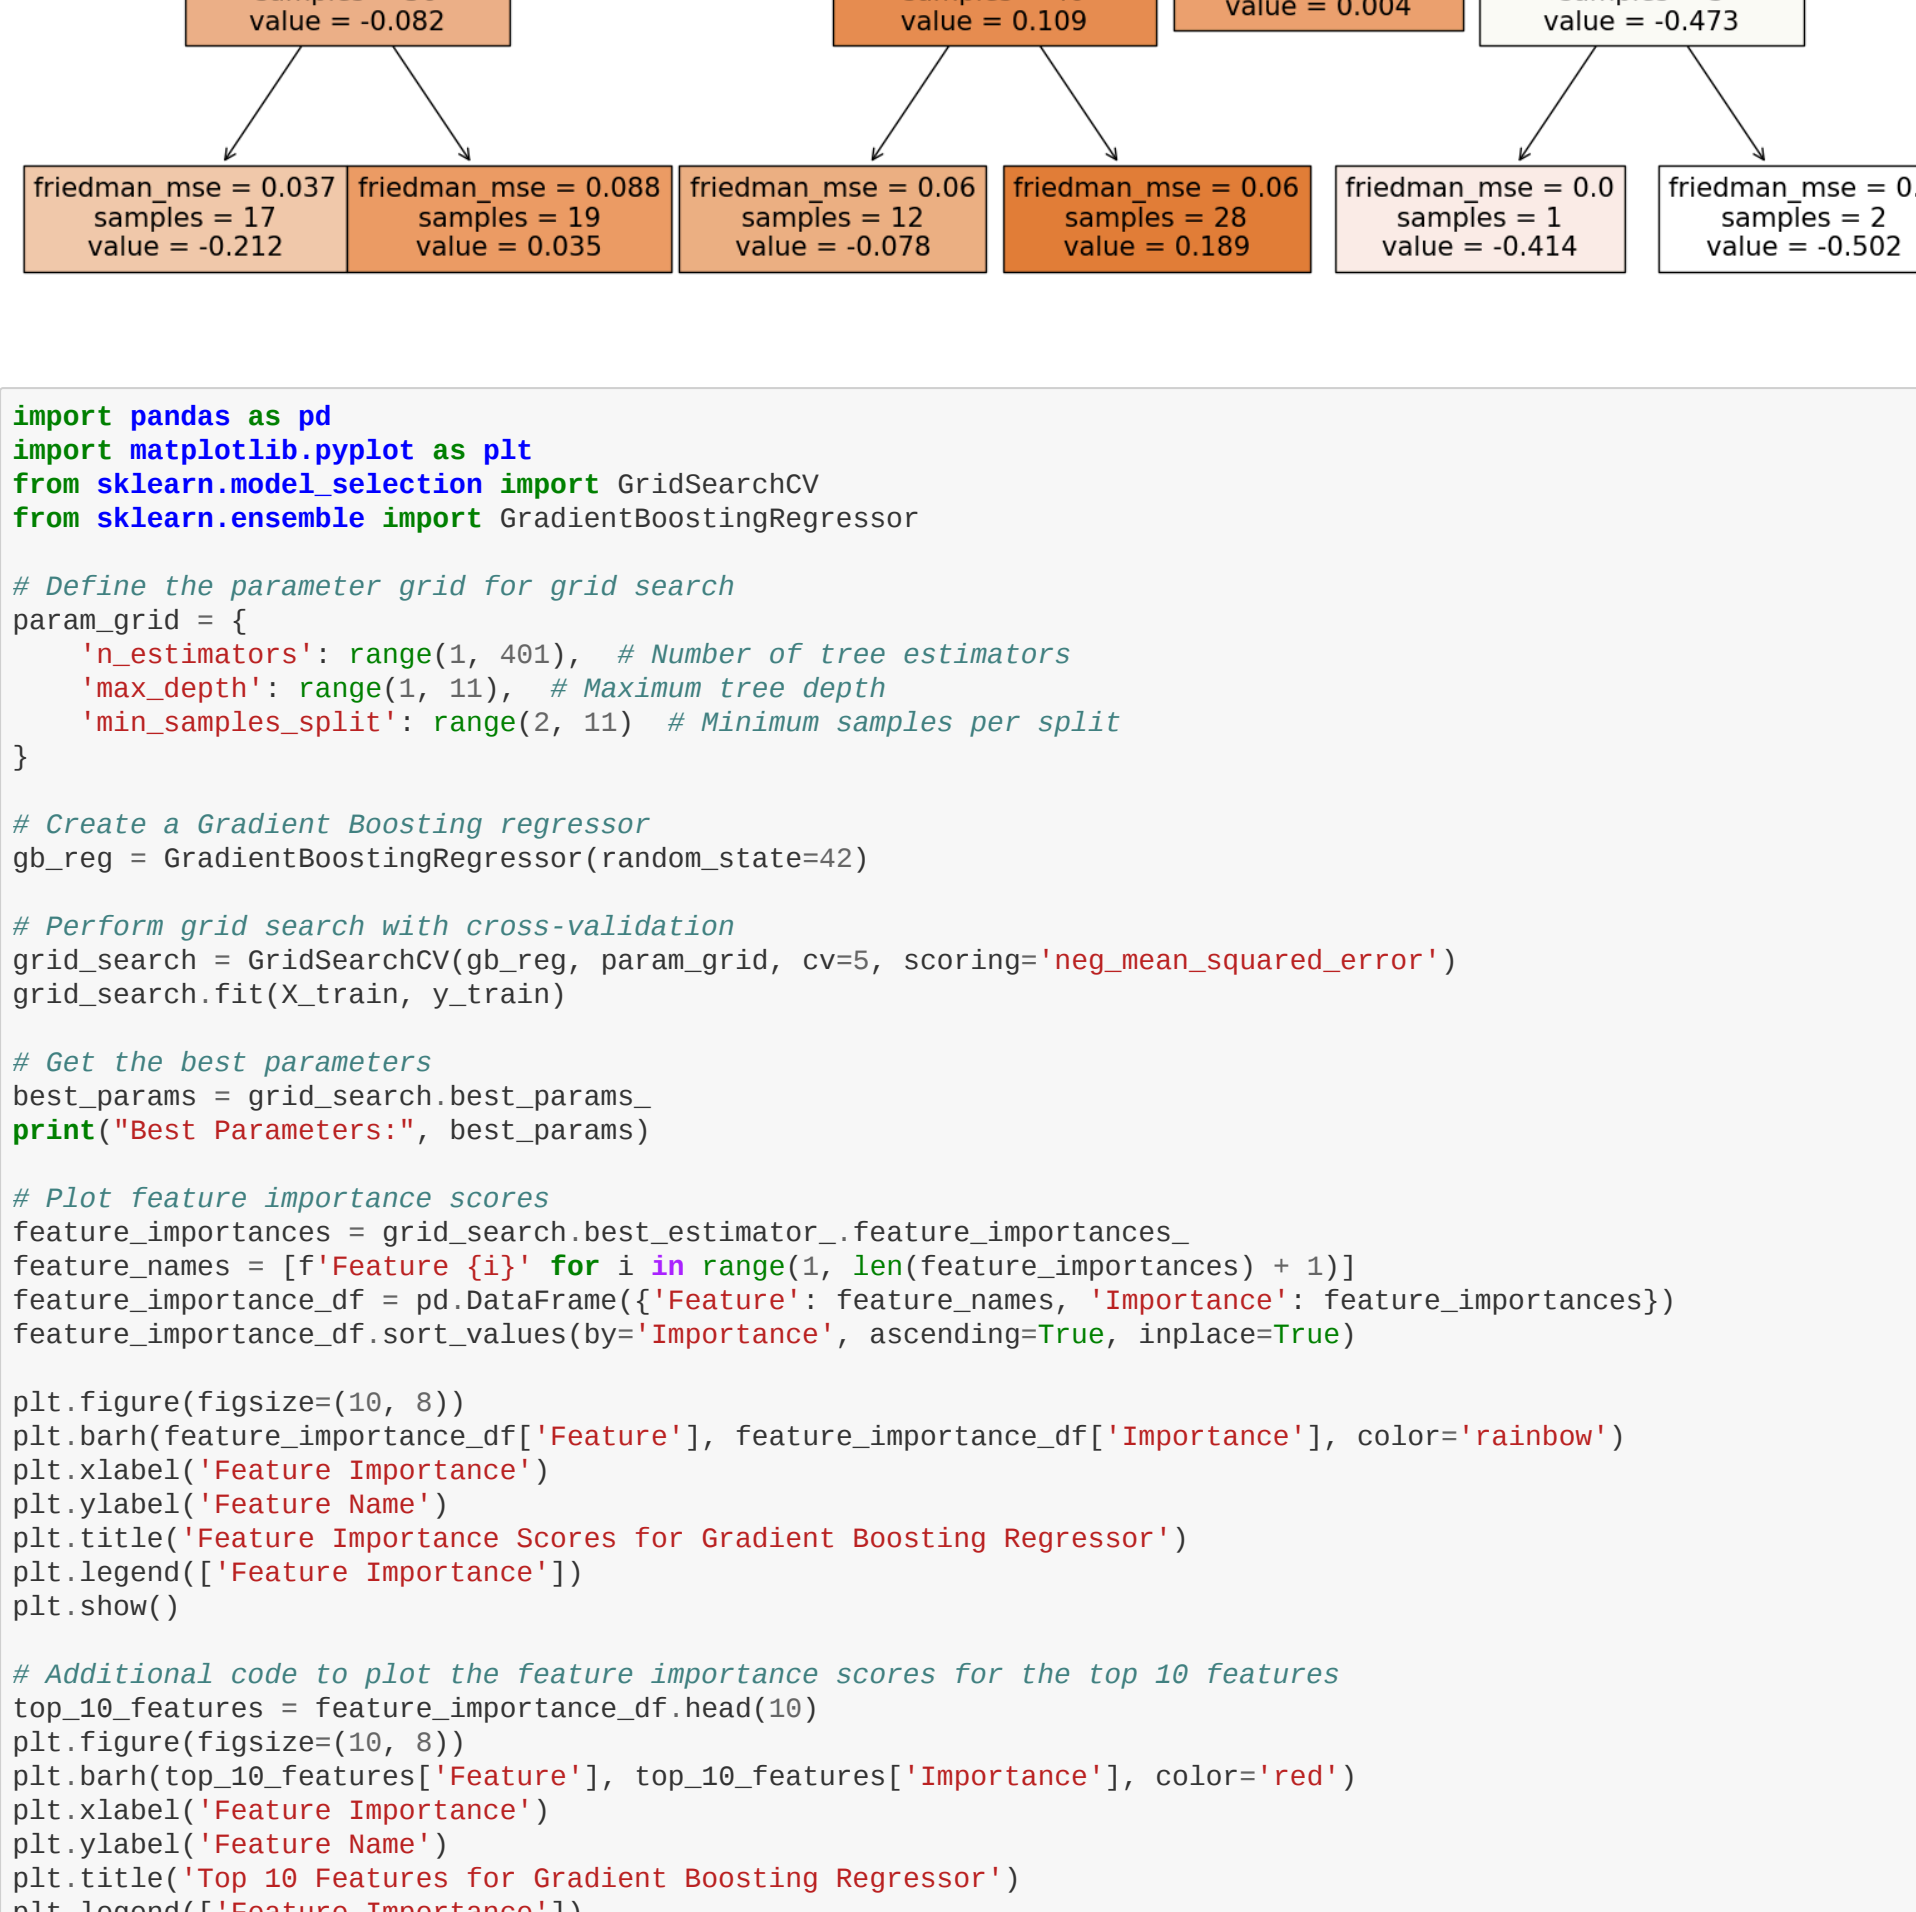
# Plot the first tree in the ensemble
```

plt.figure(figsize=(20, 10))
plot_tree(bst_reg.estimators_[0][0], filled=True) # Assessing the first tree in the first boosting iteration
plt.show()

Test MSE with Gradient Boosting: 0.10803165705663882

```
graph TD;
    Root["x[8] <= 0.966  
friedman_mse = 0.091  
samples = 80  
value = 0.0"]
    Left["x[8] <= 0.499  
friedman_mse = 0.086  
samples = 76  
value = 0.019"]
    Right["x[5] <= 0.19  
friedman_mse = 0.044  
samples = 4  
value = -0.353"]
    LeftLeft["x[4] <= 0.588  
friedman_mse = 0.079  
samples = 36"]
    LeftRight["x[5] <= 0.357  
friedman_mse = 0.075  
samples = 40"]
    LeftLeftLeft["x[4] <= 0.588  
friedman_mse = 0.079  
samples = 36"]
    LeftLeftRight["x[4] > 0.588  
friedman_mse = 0.075  
samples = 40"]
    LeftRightLeft["x[5] <= 0.0  
friedman_mse = 0.0  
samples = 1"]
    LeftRightRight["x[4] <= 0.437  
friedman_mse = 0.002  
samples = 3"]

    Root --> Left
    Root --> Right
    Left --> LeftLeft
    Left --> LeftRight
    LeftLeft --> LeftLeftLeft
    LeftLeft --> LeftLeftRight
    LeftRight --> LeftRightLeft
    LeftRight --> LeftRightRight
```



```
In [ ]: # Feature Importance }
plt.show()

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.tree import DecisionTreeClassifier

# Load the dataset
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/spambase/spambase.data"
spam_data = pd.read_csv(url, header=None)

# Split features and target variable
X = spam_data.iloc[:, :-1].values
y = spam_data.iloc[:, -1].values

# Split the dataset into training and testing sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create a bootstrap sample of size 1000
bootstrap_sample_size = 1000
bootstrap_indices = np.random.choice(len(X_train), size=bootstrap_sample_size, replace=True)

# Define a function to train decision tree classifiers with different values of p
```

```
def train_decision_tree(X_train, y_train, p):
    selected_features_indices = random.choice(X_train.shape[1], size=p, replace=False)
    X_train_selected = X_train[:, selected_features_indices]
    dt_classifier = DecisionTreeClassifier(max_depth=0)
    dt_classifier.fit(X_train_selected, y_train)
    return dt_classifier

# Define a function to calculate cross-validation error
def calculate_cross_val_error(dt_classifier, X_train, y_train):
    cv_scores = cross_val_score(dt_classifier, X_train, y_train, cv=5, scoring='accuracy')
    return 1 - np.mean(cv_scores)

# Vary the value of p and find the p that results in the lowest cross-validation error
best_p = None
lowest_cv_error = float('inf')
for p in range(1, X_train.shape[1] + 1):
    dt_classifier = train_decision_tree(X_train[bootstrap_indices], y_train[bootstrap_indices], p)
    cv_error = calculate_cross_val_error(dt_classifier, X_train, y_train)
    if cv_error < lowest_cv_error:
        lowest_cv_error = cv_error
        best_p = p

# Train a decision tree classifier with the best value of p
best_dt_classifier = train_decision_tree(X_train[bootstrap_indices], y_train[bootstrap_indices], best_p)
```

```
# Print the best value of p and the decision tree classifier
print("Best value of p: ", best_p)
print("Decision Tree Classifier with max depth 6 trained using the bootstrap sample and selected features:", best_dt_classifier)

Best value of p: 14
Decision Tree Classifier with max depth 6 trained using the bootstrap sample and selected features: DecisionTreeClassifier(max_depth=6)

In [ ]: import numpy as np
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import f1_score, roc_auc_score
from sklearn.model_selection import train_test_split

# Load the dataset:
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/spambase/spambase.data"
data = pd.read_csv(url, header=None)

# Split the data into features and target variable
x = data.iloc[:, :-1] # Features
y = data.iloc[:, -1] # Target variable
```

```
# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Define the range of values for p and T
p_values = range(1, X.shape[1] + 1) # All features
T_values = [1, 50, 100, 150, 200, 300, 400]

# Define a function to train a decision tree classifier with specified parameters
def train_decision_tree(p, T):
    dt_classifier = DecisionTreeClassifier(max_depth=max_depth)
    dt_classifier.fit(X, y)
    return dt_classifier

# Define a function to combine predictions from multiple decision tree classifiers
def combine_predictions(classifiers, X):
    predictions = np.zeros(X.shape[0], len(classifiers))
    for i, classifier in enumerate(classifiers):
        predictions[:, i] = classifier.predict(X)
    combined_predictions = np.mean(predictions, axis=1)
    return combined_predictions

# Define a function to evaluate combined predictions
def evaluate_combined_predictions(true, y_pred):
    # Round predictions to the nearest integer
```

```

y_pred_rounded = np.round(y_pred)
# Calculate F1 score
f1 = f1_score(y_true, y_pred_rounded)
# Calculate AUC score
auc = roc_auc_score(y_true, y_pred)
# Calculate training error
training_error = np.mean(y_true - y_pred) ** 2
return training_error, f1, auc

# Initialize lists to store results
best_p_values = []
training_errors = []
f1_scores = []
auc_scores = []

# Iterate over each value of p
for best_p in p_values:
    classifiers = []
    for i in range(100):
        bootstrap_indices = np.random.choice(len(X_train), size=len(X_train), replace=True)
        y_bootstrap = X_train.iloc[bootstrap_indices] # Use list to store rows by index
        y_bootstrap = y_train.iloc[bootstrap_indices]
        classifier = train_decision_tree(X_bootstrap, y_bootstrap, c) # Set max_depth to 6
        classifiers.append(c, classifier)

```

```

combined_predictions_train = combine_predictions(classifiers, X_train)
training_auc, f1 = evaluate_combined_predictions(p_train, combined_predictions_train)
best_p_values.append(best_p)
training_errors.append(training_error)
f1_scores.append(f1)
auc_scores.append(auc)

# Print the results
for p, training_error, f1, auc in zip(best_p_values, training_errors, f1_scores, auc_scores):
    print("P({p}): {p}, Training Error:{training_error:.4f}, F1 Score:{f1:.4f}, AUC:{auc:.4f}")

In [ ]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, f1_score, roc_auc_score

# Initialize list of number of trees
num_trees_list = [10, 50, 100]

# Train and evaluate Random Forest models with different number of trees
for num_trees in num_trees_list:
    # Create Random Forest classifier
    rf_classifier = RandomForestClassifier(n_estimators=num_trees, random_state=42)

    # Train the classifier
    rf_classifier.fit(X_train, y_train)

```

```
# Predictions on training set
y_train_pred = rf_classifier.predict(X_train)

# Metrics on training set
train_accuracy = accuracy_score(y_train, y_train_pred)
train_f1 = f1_score(y_train, y_train_pred)
train_auc = roc_auc_score(y_train, y_train_pred)

# Predictions on testing set
y_test_pred = rf_classifier.predict(X_test)

# Metrics on testing set
test_accuracy = accuracy_score(y_test, y_test_pred)
test_f1 = f1_score(y_test, y_test_pred)
test_auc = roc_auc_score(y_test, y_test_pred)

# Print metrics
print(f"Number of Trees: {num_trees}")
print(f"Training Set Metrics:")
print(f"Accuracy: {train_accuracy:.4f}, F1 Score: {train_f1:.4f}, AUC: {train_auc:.4f}")
print(f"Testing Set Metrics:")
print(f"Accuracy: {test_accuracy:.4f}, F1 Score: {test_f1:.4f}, AUC: {test_auc:.4f}")

# Feature Importances
feature_importances = rf_classifier.feature_importances_
```

```
# Get indices of top 10 features
top_10_indices = Feature_Importance.argsort()[1:-10][::-1]
# Print top 10 features
print("Top 10 Features:")
for i, idx in enumerate(top_10_indices):
    print(f"({i} + 1). Feature ({idx}): Importance = {Feature_Importance[idx]:.4f}")
print("\n")

Number of Trees: 10
Training Set Metrics:
Accuracy: 0.9846, F1 Score: 0.9929, AUC: 0.9934
Testing Set Metrics:
Accuracy: 0.9403, F1 Score: 0.9262, AUC: 0.9329
Top 10 features:
1. Feature 6: Importance = 0.1209
2. Feature 51: Importance = 0.2837
3. Feature 55: Importance = 0.0880
4. Feature 52: Importance = 0.0794
5. Feature 23: Importance = 0.0620
6. Feature 15: Importance = 0.0546
7. Feature 54: Importance = 0.0554
8. Feature 56: Importance = 0.0441
9. Feature 24: Importance = 0.0384
10. Feature 22: Importance = 0.0295
```

```

Number of Trees: 50
Training Set Metrics:
Accuracy: 0.9995, F1 Score: 0.9993, AUC: 0.9994
Testing Set Metrics:
Accuracy: 0.9498, F1 Score: 0.9377, AUC: 0.9435
Top 10 Features:
1. Feature 53: Importance = 0.1143
2. Feature 52: Importance = 0.1851
3. Feature 6: Importance = 0.0894
4. Feature 55: Importance = 0.0685
5. Feature 15: Importance = 0.0634
6. Feature 56: Importance = 0.0517
7. Feature 54: Importance = 0.0503
8. Feature 24: Importance = 0.0425
9. Feature 20: Importance = 0.0350
10. Feature 18: Importance = 0.0334

Number of Trees: 100
Training Set Metrics:
Accuracy: 0.9995, F1 Score: 0.9993, AUC: 0.9993
Testing Set Metrics:

```

Accuracy: 0.9555, F1 Score: 0.9458, AUC: 0.9505
Top 10 Features:

1. Feature 51: Importance = 0.1138
2. Feature 52: Importance = 0.0969
3. Feature 6: Importance = 0.0819
4. Feature 15: Importance = 0.0671
5. Feature 55: Importance = 0.0680
6. Feature 54: Importance = 0.0579
7. Feature 56: Importance = 0.0524
8. Feature 20: Importance = 0.0463
9. Feature 24: Importance = 0.0424
10. Feature 18: Importance = 0.0329